

Using the Foundations and Trends[®] L^AT_EX Class

Instructions for Creating an FnT Article

Suggested Citation: Alet Heezemans and Mike Casey (2018), "Using the Foundations and Trends[®] L^AT_EX Class", : Vol. xx, No. xx, pp 1–18. DOI: 10.1561/XXXXXXXXXX.

Alet Heezemans
now publishers, Inc.
alet.heezemans@nowpublishers.com

Mike Casey
now publishers, Inc.
mike.casey@nowpublishers.com

This article may be used only for the purpose of research, teaching, and/or private study. Commercial use or systematic downloading (by robots or other automatic processes) is prohibited without explicit Publisher approval.

now
the essence of knowledge
Boston — Delft

Contents

1	Introduction to Dynamic Programming	2
1.1	Dynamic Programming Setting	2
1.2	Dynamic Programming Solving	8
2	Dynamic Programming for Deterministic System	13
2.1	Deterministic Finite-State Problem	13
2.2	Generic Stochastic Path Problems	15
2.3	General Shortest Path Algorithms	17
3	Dynamic Programming with perfect information	21
3.1	Linear-Quadratic Problems	21
3.2	Inventory Control	25
3.3	Optimal Stopping Problem	28
4	Dynamic Programming with imperfect information	32
4.1	Problem Setting	32
4.2	Reformulation as a perfect state information system	33
4.3	Machine Repair Example	34
5	The Distribution and Installation	36
5.1	Pre-requisites	36
5.2	The Distribution	36

5.3	Installation	37
6	Quick Start	39
6.1	<code>\documentclass</code>	39
6.2	<code>\issuesetup</code>	40
6.3	<code>\maintitleauthorlist</code>	40
6.4	<code>\author</code> and <code>\affil</code>	40
6.5	<code>\addbibresource</code>	40
7	Style Guidelines and \LaTeX Conventions	41
7.1	Abstract	41
7.2	Acknowledgements	41
7.3	References	41
7.4	Citations	42
7.5	Preface and Other Special Chapters	42
7.6	Long Chapter and Section Names	43
7.7	Internet Addresses	43
8	Compiling Your $\text{\textsf{FnT}}$ Article	44
8.1	Compiling Your Article Prior to Submission	44
8.2	Preparing the Final Versions	45
8.3	Compiling The Final Versions	45
8.4	Wednesday	46
	Acknowledgements	51
	Appendices	52
A	Journal Codes	53
B	Files Produced During Compilation	55
	References	56

Using the Foundations and Trends[®] L^AT_EX Class

Alet Heezemans¹ and Mike Casey²

¹*now publishers, Inc.; alet.heezemans@nowpublishers.com*

²*now publishers, Inc.; mike.casey@nowpublishers.com*

ABSTRACT

This document describes how to prepare a Foundations and Trends[®] article in L^AT_EX . The accompanying L^AT_EX source file FnTarticle.tex (that produces this output) is an example of such a file.

1

Introduction to Dynamic Programming

1.1 Dynamic Programming Setting

Consider a *generic* optimization problem

$$\min_{u \in U} g(u)$$

where u is the decision variable, $g(u)$ is the cost function, and U is the constraint set. There are several categories of problems:

- Discrete (i.e., U is finite), or continuous
- Linear (i.e., g is linear and U is polyhedral) or nonlinear
- Stochastic or deterministic: For stochastic problems, the cost involves a stochastic parameter ω , which is averaged:

$$g(u) = \mathbb{E}_{\omega}\{G(u, \omega)\}$$

where ω is a random parameter.

The dynamic programming (dp) is able to deal with complex stochastic problems where ω becomes available in stages, and the decisions are also made in stages based on the information.

1.1.1 Basic Structure of Stochastic DP

The DP usually involves the following elements:

- Discrete-time system:

$$x_{k+1} = f_k(x_k, u_k, \omega_k), \quad k = 0, 1, \dots, N-1$$

where

- k denotes the stage / discrete time
 - x_k denotes the state for stage k , which *summarizes past information* that is relevant for future decision
 - u_k denotes the control variable, i.e., a decision to be selected at time k from a given set. When making the decision, the state x_k is known.
 - w_k denotes the *random parameter / disturbance / noise* at stage k .
 - N is the *Horizon* or the number of times that the control is applied.
- A cost function that is *addictive* over time:

$$\mathbb{E} \left\{ g_N(x_N) + \sum_{k=0}^{N-1} g_k(x_k, u_k, \omega_k) \right\}$$

Remark 1.1. We may use the transition matrix as the alternative *discrete-system* description, i.e., suppose the transition matrix $P(x_{k+1} \mid x_k, u_k)$ is available, then

$$x_{k+1} = \omega_k, \text{ with } P(\omega_k \mid x_k, u_k) = P(x_{k+1} \mid x_k, u_k)$$

Several assumptions are made for common dp problem:

- The set of possible values for control u_k depend at most on x_k but not on *prior* x or u
- The probability distribution of ω_k does not depend on past values $\omega_{k-1}, \dots, \omega_0$, i.e., $\omega_k \sim P(\cdot; x_k, u_k)$, but may depend on x_k and u_k , since otherwise past values of ω would be useful for future optimization by applying *state augmentation*.




Figure 1.1: Inventory Model

1.1.2 Inventory Control Example

The inventory model is summarized in the figure below: In this case, the discrete-time system is

$$x_{k+1} = f_k(x_k, u_k, \omega_k) = x_k + u_k - \omega_k$$

with the cost function that is *additive* over time:

$$\mathbb{E} \left\{ g_N(x_N) + \sum_{k=0}^{N-1} g_k(x_k, u_k, \omega_k) \right\} = \mathbb{E} \left\{ \sum_{k=0}^{N-1} (c \cdot u_k + r \cdot (x_k + u_k - \omega_k)) \right\}$$

Our goal is to optimize the cost function over *policies*, and the policies $u_k = \mu_k(x_k)$ maps states into controls. The discrete-time system for this problem is deterministic, but the cost function is a expected value for a stochastic rv.

1.1.3 Deterministic Finite-State Problem

Consider a deterministic dp problem. We aim to find the optimal sequence of four operations A, B, C, D , where A must *precede* B and C must *precede* D . Assume the startup cost to be S_A, S_C , and the setup transition cost from operation m to operation n is C_{mn} .

We may list all possible scheduling in the figure below:

The discrete-time system for this problem is deterministic, and the cost function is deterministic as well.




Figure 1.2: Scheduling Problem

1.1.4 Stochastic Finite-State Problem

The goal is to find the two-game chess match strategy.

- The *timid* play draws with probability $p_d > 0$ and loses with probability $1 - p_d$
- The *bold* play wins with probability $p_w < \frac{1}{2}$ and loses with probability $1 - p_w$.


We may list all possible strategies result in the figure (1.3). The discrete-time system for this problem is stochastic, and the cost function is stochastic as well.

1.1.5 Formal Setting of DP

The formal setting of DP is as follows:

- The system transition satisfies

$$x_{k+1} = f(x_k, u_k, \omega_k), \quad k = 0, 1, \dots, N - 1$$

**Figure 1.3:** Two-Game Chess Match Strategies

- The control constraints only depends on x_k , i.e., $u_k \in U_k(x_k)$
- The probability distribution of ω_k is $P_k(\cdot \mid x_k, u_k)$
- The policy is a collection of functions $\pi = \{\mu_0, \dots, \mu_{N-1}\}$, where μ_k maps states x_k into controls $u_k = \mu_k(x_k) \in U_k(x_k)$ for all x_k
- The expected cost of π starting at x_0 is given by

$$J_\pi(x_0) = \mathbb{E} \left\{ g_N(x_N) + \sum_{k=0}^{N-1} g_k(x_k, \mu_k(x_k), \omega_k) \right\}$$

The goal is to minimize the cost function:

$$J^*(x_0) = \min_{\pi} J_\pi(x_0),$$

where the optimal policy π^* satisfies

$$J_{\pi^*}(x_0) = J^*(x_0).$$

Note that π^* is independent of x_0 .

1.1.6 Significance of Feedback

The closed-loop policies refer to that the controller can adapt to the unexpected values of the state, while the open-loop policies cannot get access to this information.

- For deterministic problems, the *open-loop policies* is as good as the *closed-loop policies*
- For stochastic problems, the reduction in the cost gained by closed-loop is called the *value of information*.

For example, consider the chess match problem,

1. The probability of win using optimal open-loop policy is

$$p_w^2 + p_w(1 - p_w) \cdot \max(2p_w, p_d)$$

2. The probability of win using optimal closed-loop policy is

$$p_w^2(2 - p_2) + p_w(1 - p_w) \cdot p_d$$

3. Therefore, we see that the value of information is sometimes positive:

$$\text{Value of Information} = p_w(1 - p_w) \min(p_w, p_d - p_w)$$

Variants of DP In the second-half of this course, we may consider several variants of DP problems:

- Continuous-time problems
- Imperfect state information problems
- Infinite horizon problems
- Suboptimal Control

1.2 Dynamic Programming Solving

1.2.1 Principle of Optimality

Suppose $\pi^* = \{\mu_0^*, \dots, \mu_{N-1}^*\}$ be optimal policy. Consider the *tail subproblem*, i.e., we are at x_i at time i and wish to minimize the *cost-to-go* from time i to time N :

$$\mathbb{E} \left\{ g_N(x_N) + \sum_{k=i}^{N-1} g_k(x_k, \mu_k(x_k), \omega_k) \right\}$$

Consider the *tail policy* $\{\mu_i^*, \mu_{i+1}^*, \dots, \mu_{N-1}^*\}$.

Theorem 1.1 (Principle of Optimality). The tail policy is optimal for the tail subproblem, i.e., the optimization of the future does not depend on what we did in the past.

Remark 1.2. Consider an auto travel analogy. Suppose that the fastest route from A to C passes the place B. The principle of optimality highlights that *the B to C portion of the route is also the fastest route of the trip that starts from B and ends at C*.

The DP aims to solve all tail subproblems of N horizons. The DP algorithm is based on the idea as follows: It proceeds sequentially, by solving all the tail subproblems of a given time length, based on the solution of the tail subproblems of shorter time length.

1.2.2 DP algorithm for scheduling problem

Consider the scheduling problem discussed in subsection (1.1.3). The startup cost and transition cost is summarized in the figure (1.4).




Figure 1.4: Scheduling Problem Cost Description

Tail Problem of Length 1: The subproblems involving one operations are trivial to solve. State $ABC, ACB, ACD, CAB, CAD, CDA$ correspond to the tail cost 6, 1, 3, 1, 3, 2.

Tail Problem of Length 2: Solving the subproblems involving two operations is based on the tail problem of length 1. State AB, AC, CA, CD corresponds to the tail cost 9, $\min(4+1, 6+3) = 5$, $\min(2+1, 4+3) = 3$, 5.

Tail Problem of Length 3: Solving the subproblems involving three operations is based on the tail problem of length 2. State A, C corresponds to the tail cost $\min(2+9, 3+5) = 8$, $\min(4+3, 6+5) = 7$.

Tail Problem of Length 4: Solving the subproblems involving four operations is based on the tail problem of length 3. The initial state corresponds to the tail cost $\min(5+8, 3+7) = 10$.

After computing the optimal cost, we construct the optimal solution by starting at the initial node and proceeding forward, where each time we choose the operation that starts the *optimal* schedual for the

corresponding *tail subproblem*. Therefore, at each state-time pair, we need to record both the optimal cost-to-go and the optimal decision.

1.2.3 Stochastic Inventory Example

Consider the inventory problem in subsection (1.1.2). The tail subproblems of length 1 aims to solve

$$J_{N-1}(x_{N-1}) = \min_{u_{N-1} \geq 0} \mathbb{E}_{\omega_{N-1}} \{c \cdot u_{N-1} + r \cdot (x_{N-1} + u_{N-1} - \omega_{N-1})\}, \forall x_{N-1}$$

The tail subproblems of length $N - k$ aims to solve

$$J_k(x_k) = \min_{u_k \geq 0} \mathbb{E}_{\omega_k} \{c \cdot u_k + r \cdot (x_k + u_k - \omega_k) + J_{k+1}(x_k + u_k - \omega_k)\}, \forall x_k$$

Here $J_0(x_0)$ is the optimal cost with initial state x_0 .

1.2.4 Formal Dynamic Programming Algorithm

The dynamic programming aims to find the policy by starting with

$$J_N(x_N) = g_N(x_N),$$

and go backwards in time from period $N - 1$ to period 0:


$$J_k(x_k) = \min_{u_k \in U_k(x_k)} \mathbb{E}_{\omega_k} \{g_k(x_k, u_k, \omega_k) + J_{k+1}(f_k(x_k, u_k, \omega_k))\}, k = 0, \dots, N-1. \quad (1.1)$$

In such case, $J_0(x_0)$ is generated at the last step, and is equal to the optimal cost $J^*(x_0)$. Furthermore, if $u_k^* = \mu_k^*(x_k)$ minimizes the RHS of (1.1) for each x_k and k , then $\pi^* = \{\mu_0^*, \dots, \mu_{N-1}^*\}$ is optimal.

Proof. Proof by induction that $J_k(x_k)$ is equal to $J_k^*(x_k)$, which is defined as the optimal cost of the tail subproblem that starts at time k at state x_k . \square

1.2.5 Linear-Quadratic Analytic Example

Consider a certain material is passed through a sequence of two ovens shown in Figure (1.5).

**Figure 1.5:** Linear-Quadratic Example

- The system is given by

$$x_{k+1} = (1 - a)x_k + au_k, \quad k = 0, 1,$$

where a is a given scalar from the interval $(0, 1)$.

- Define the cost function

$$r(x_2 - T)^2 + u_0^2 + u_1^2$$

where r is a given positive scalar.

- We may apply the dp algorithm as follows:

$$J_2(x_2) = r(x_2 - T)^2$$

$$J_1(x_1) = \min_{u_1} \left[u_1^2 + r((1 - a)x_1 + au_1 - T)^2 \right]$$

$$J_0(x_0) = \min_{u_0} \left[u_0^2 + J_1((1 - a)x_0 + au_0) \right]$$

1.2.6 State Augmentation

When the assumptions of the basic problem are violated, then we need to consider reformulating the state. The DP algorithm still applies, but the problem gets bigger.

Example Consider the system becomes

$$x_{k+1} = f_k(x_k, x_{k-1}, u_k, \omega_k)$$

Introducing the additional state variable $y_k = x_{k-1}$ and view $\tilde{x}_k = (x_k, y_k)$. Therefore, the new system takes the form

$$\begin{pmatrix} x_{k+1} \\ y_{k+1} \end{pmatrix} = \begin{pmatrix} f_k(x_k, y_k, u_k, \omega_k) \\ x_k \end{pmatrix}$$

The DP algorithm applies to the reformulated problem:

$$J_k(x_k, x_{k-1}) = \min_{u_k \in U_k(x_k)} \mathbb{E}_{\omega_k} \{g_k(x_k, u_k, \omega_k) + J_{k+1}(f_k(x_k, x_{k-1}, u_k, \omega_k), x_k)\}$$

2

Dynamic Programming for Deterministic System

2.1 Deterministic Finite-State Problem

Consider a shortest path problem shown in Figure (2.1). We can convert the problem into a dynamic programming system as follows:

- The states correspond to the nodes in each stage.
- The controls correspond to the acrs in each stage.
- The control sequences for this open-loop system correspond to




Figure 2.1: Illustration of shortest path problem

paths from initial state to the terminal states.

- Let a_{ij}^k denote the cost of transition (length of the arc) from the state $i \in S_k$ to state $j \in S_{k+1}$ at time k
- The cost of the control sequence is the cost of the corresponding path (total length of the path)

DP algorithms The backward algorithm applies as follows:

$$J_N(i) = a_{ii}^N, \quad i \in S_N$$

$$J_k(i) = \min_{j \in S_{k+1}} [a_{ij}^k + J_{k+1}(j)], \quad i \in S_k, k = N-1, N-2, \dots, 0.$$

Here the length of the shortest path from s to t equals to the $J_0(s)$.

Note that the optimal path $s \rightarrow t$ is also an optimal path $t \rightarrow s$ for a *reversed* shortest path problem, where the direction of each arc is reversed with its length unchanged.

Therefore, we can apply the backward DP algorithm for the reversed problem, which is so called the *forward* DP algorithm:

$$\tilde{J}_N(j) = a_{sj}^0, \quad j \in S_1,$$

$$\tilde{J}_k(j) = \min_{i \in S_{N-k}} [a_{ij}^{N-k} + \tilde{J}_{k+1}(i)], \quad j \in S_{N-k+1}$$

The optimal cost for the whole problem is $\tilde{J}_0(t) = \min_{i \in S_N} [a_{it}^N + \tilde{J}_1(i)]$. The interpretation for this forward DP algorithm is that we can view $\tilde{J}_k(j)$ as the *optimal cost-to-arrive* to state j from initial state s .

Remark 2.1. There is no *forward* DP algorithm for *stochastic problems*. This is because that we cannot restrict ourselves to open-loop sequences for the stochastic problems. In other words, due to the uncertainty, the concept of *optimal-cost-to-arrive* at the state x_k does not make sense, i.e., it's impossible to guarantee (with probability 1) that any given state can be reached.

Fortunately, even for stochastic problems, the concept of *optimal cost-to-go* from any state x_k makes clear sense.

2.2 Generic Stochastic Path Problems

Consider first a general deterministic shortest path(s) problems as follows:

- Let $\{1, \dots, N, t\}$ denote the nodes of a graph, where t is the destination.
- The a_{ij} denotes the cost of moving from node i to node j
- Our goal is to find a shortest (minimum cost) path from each node i to node t
- Assumption: *All cycles have non-negative length*, i.e., the optimal path *need not take more than N moves*.
- We can reformulate this problem as one where we require *exactly* N moves by allowing *degerate moves* from a node i to itself with cost $a_{ii} = 0$. Moreover, we let

$$J_k(i) = \text{optimal cost of getting from } i \text{ to } t \text{ in } N - k \text{ moves}$$

$$J_0(i) = \text{cost of the optimal path from } i \text{ to } t$$

- Therefore, we apply the DP algorithm as follows:

$$J_k(i) = \min_{j=1, \dots, N} [a_{ij} + J_{k+1}(j)], \quad k = 0, 1, \dots, N - 2,$$

$$\text{with } J_{N-1}(i) = a_{it}, i = 1, \dots, N.$$

For example, consider the shortest paths problem shown in Fig (2.2). Assume the destination node is the node 5. The state i means the current position is in the node i . In such case, note that

$$J_{N-1}(i) = a_{it}, \quad i = 1, \dots, N$$

$$J_k(i) = \min_{j=1, \dots, N} [a_{ij} + J_{k+1}(j)], \quad k = 0, 1, \dots, N - 2$$

The solution for this problem is shown in Fig. (2.2b), where the arrow indicate the optimal move to node 5 from the position at current stage.




Figure 2.2: One Example of shortest paths problems

2.2.1 HMM setting and solving

Now consider the *Hidden Markov Models*. Suppose that the Markov chain is with transition probability p_{ij} , but the state transitions are hidden from our view. For each transition, we obtain an *observation* instead. Let $r(z; i, j)$ denote the probability of observation taking value z given the transition $i \rightarrow j$. The *trajectory estimation problem* is that, given the observation sequence $Z_N = \{z_1, \dots, z_N\}$, what's the *most likely* state transition sequence $\hat{X}_N := \{\hat{x}_0, \hat{x}_1, \dots, \hat{x}_N\}$, i.e., the one equals to

$$\arg \min_{\text{all } X_N} p(X_N | Z_N) = \frac{p(X_N, Z_N)}{p(Z_N)}$$

We can rewrite the probability $p(X_N, Z_N)$ as follows:

$$\begin{aligned} p(X_N, Z_N) &:= p(x_{0:N}, z_{1:N}) \\ &= \pi_{x_0} \cdot p(x_{1:N}, z_{1:N} | x_0) \\ &= \pi_{x_0} p(x_1, z_1 | x_0) p(x_{2:N}, z_{2:N} | x_{0:1}, z_1) \\ &= \pi_{x_0} p_{x_0 \rightarrow x_1} r(z_1; x_0, x_1) p(x_{2:N}, z_{2:N} | x_{0:1}, z_1) \end{aligned}$$

Following the similar idea, we compute

$$\begin{aligned} p(x_{2:N}, z_{2:N} | x_{0:1}, z_1) &= p(x_2, z_2 | x_{0:1}, z_1) p(x_{3:N}, z_{3:N} | x_{0:2}, z_{1:2}) \\ &= p_{x_1 \rightarrow x_2} r(z_2; x_1, x_2) p(x_{3:N}, z_{3:N} | x_{0:2}, z_{1:2}) \end{aligned}$$

Therefore, we derive

$$p(X_N, Z_N) = \pi_{x_0} \prod_{k=1}^N p_{x_{k-1} \rightarrow x_k} r(z_k; x_{k-1}, x_k)$$

Thus the problem is equivalent to

$$\text{minimize} \quad -\ln(\pi_{x_0}) - \sum_{k=1}^N \ln(p_{x_{k-1} \rightarrow x_k} r(z_k; x_{k-1}, x_k))$$

over all possible sequences $\{x_0, x_1, \dots, x_N\}$. Therefore, due to the structure of our objective function, we can apply a distributed algorithm (e.g., dynamic programming) to solve it.

2.3 General Shortest Path Algorithms

Next, we discuss some non-DP shortest path algorithms. Note that they all can *only* be used to solve deterministic finite-state problems. The advantage is that they avoid computing the optimal cost-to-go for *every* state. This is essential for problems with *huge* state spaces. (see examples in combinatorial optimization)

2.3.1 Label Correcting Methods

The idea of this kind of algorithm is to *progressively* discover shorter paths from the origin to *any other* node i , and to maintain the length of the shortest path found *so far* in a variable d_i (called the *label* of i).

- If a shorter path to i is found, then d_i is reduced, and the algorithm checks if the labels d_j can be reduced, where j is the children of i , i.e., whether labels d_j 's can be reduced by setting them into $d_i + a_{ij}$.
- This algorithm makes use of a list of nodes called *OPEN* (or *candidate list*). This list contains nodes that are currently active such that they are candidates for further examination by the algorithm, and possible inclusion in the shortest path.

Initially, OPEN contains only s . Each node that has entered the OPEN at least once (except s) is assigned a *parent* (refers to some node). The parent nodes are not necessary for the computation of shortest distance, but for tracing the shortest path to the origin after the algorithm terminates.

Notations

- d_i denotes the *label* of i , i.e., the length of the shortest path from s to i found so far. Initially $d_s = 0, d_i = \infty, \forall i \neq s$
- UPPER: The label d_t to the destination
- OPEN list: contains nodes that are currently active in the sense that they are candidates for further examination (initially OPEN = $\{s\}$)

The label correcting methods work as follows:

Algorithm 1 The Label Correcting Methods

- 1: (Node Removal): Remove a node i from OPEN and for each child j of i , do step 2
 - 2: (**Node Insertion Test**): If $d_i + a_{ij} < \min\{d_j, \text{UPPER}\}$, set $d_j = d_i + a_{ij}$, and set i to be the parent of j . In addition, if $j \neq t$, place j into the OPEN if $j \notin \text{OPEN}$; otherwise set UPPER to the new value $d_i + a_{it}$ of d_t
 - 3: (**Termination Test**): If OPEN is empty, terminate, else go to step 1.
-

Here we show how to apply the label correcting algorithm into the travelling salesman problem shown below.

Iteation Number	Node Exiting OPEN	OPEN after Iteration	Upper
0	—	1	∞
1	1	2, 7, 10	
2	2	3, 5, 7, 10	∞
3	3	4, 5, 7, 10	∞
4	4	5, 7, 10	43
5	5	6, 7, 10	43
6	6	7, 10	13
7	7	8, 10	13
8	8	9, 10	13
9	9	10	13
10	10	\emptyset	13




Figure 2.3: Diagram Illustration of Label Correcting Methods




Figure 2.4: One example for travelling salesman problem

Note that some nodes never entered OPEN.

Finally, we give a verification for why label correcting methods work.

Proposition 2.1. If there exists at least one path from the origin to the destination, the label correcting algorithm terminates with UPPER equal to the shortest distance from the origin to the destination.

Proof. 1. Each time a node j enters OPEN, its label is decreased and becomes equal to the length of some path from i to j

2. The number of possible distinct path lengths is finite, and therefore the number of times a node can enter is finite, and the algorithm terminates.

3. Let $\{(s, j_1, \dots, j_k, t), d^*\}$ be the pair of shortest path and the corresponding shortest distance. If $\text{UPPER} > d^*$ at the termination, UPPER will also be larger than the length of all paths $(s, j_1, \dots, j_m), m = 1, \dots, k$ throughout the algorithm, i.e., the node j_k will never enter the OPEN list with d_{j_k} equal to the shortest distance from s to j_k .

Similarly, node j_{k-1} will never enter to the OPEN list with $d_{j_{k-1}}$ equal to the shortest distance from s to j_{k-1} . Continue to j_1 by induction will obtain a contradiction.

□

3

Dynamic Programming with perfect information

3.1 Linear-Quadratic Problems

Now consider the case where the future observations is linear w.r.t. current observation and controller:

$$x_{k+1} = A_k x_k + B_k u_k + \omega_k$$

The goal is to minimize the quadratic cost

$$\mathbb{E}_{\omega_k, \forall k} \left\{ x_N^T Q_N x_N + \sum_{k=0}^{N-1} (x_k^T Q_k x_k + u_k^T R_k u_k) \right\}$$

Assumptions:

1. $Q_k \succeq 0$ and $R_k \succ 0$;
2. the noise ω_k are *independent* with *zero mean*.

We may apply the DP algorithm to solve it:

$$\begin{aligned} J_N(x_N) &= x_N^T Q_N x_N \\ J_k(x_k) &= \min_{u_k} \{ x_k^T Q_k x_k + u_k^T R_k u_k + J_{k+1}(A_k x_k + B_k u_k + \omega_k) \} \end{aligned}$$

We *claim* several key facts:

Theorem 3.1. 1. $J_k(x_k)$ is *quadratic*

2. Optimal policy $\{\mu_0^*, \dots, \mu_{N-1}^*\}$ is *linear*, i.e.,

$$\mu_k^*(x_k) = L_k x_k, \quad k = 0, \dots, N-1$$

Therefore, we may apply the similar treatment to solve for $J_k(x_k)$ with $k = N-1, N-2, \dots, 1$.

Proof. By induction it is easy to verify

$$\mu_k^*(x_k) = L_k x_k, \quad J_k(x_k) = x_k^T K_k x_k + \text{constant}$$

where

$$L_k = -(B_k^T K_{k+1} B_k + R_k)^{-1} B_k^T K_{k+1} A_k, \quad k = 0, \dots, N-1$$

$$K_N = Q_N$$

$$K_k = A_k^T (K_{k+1} - K_{k+1} B_k (B_k^T K_{k+1} B_k + R_k)^{-1} B_k^T K_{k+1}) A_k + Q_k$$

The equations above are called the *discrete-time Riccatic equation*.

Similar as DP, it starts at the terminal time N and proceeds backwards. □

Asymptotic Behavior

Theorem 3.2. Consider the *Riccatic equation*

$$K_{k+1} = A^T (K_k - K_k B (B^T K_k B + R)^{-1} B^T K_k) A + Q, \quad k = 0, 1, 2, \dots$$

with the initial K_0 be arbitrary semidefinite metrix. Suppose that (A, B) is controllable, and (A, C) is observable. As a result:

1. The Riccati equation converges:

$$\lim_{k \rightarrow \infty} K_k = K,$$

where K is the unique solution of the *algebraic Riccatic equation*

$$K = A^T (K - K B (B^T K B + R)^{-1} B^T K) A + Q$$




Figure 3.1: Graphical Proof For Scalar Systems

2. The corresponding steady-state controller

$$\mu^*(x) = Lx, \quad \text{where } L = -(B^T K B + R)^{-1} B^T K A$$

is stable, i.e., the matrix $(A + BL)$ for the closed loop system

$$x_{k+1} = (A + BL)x_k + \omega_k$$

satisfies $\lim_{k \rightarrow \infty} (A + BL)^k = 0$.

Proof for scalar systems. Consider the scalar stationary system with $A \neq 0, B \neq 0, Q > 0, R > 0$. The Riccatic equation becomes

$$P_{k+1} = A^2 \left(P_k - \frac{B^2 P_k^2}{B^2 P_k + R} \right) + Q$$

Suppose that $P_{k+1} = F(P_k)$, where

$$F(P) = \frac{A^2 R P}{B^2 P + R} + Q$$

In Fig (3.1) we find that $F(P)$ has two fixed points P^* and \tilde{P} . Note that $F(P)$ is increasing on $(-R/B^2, \infty)$. Therefore, the Riccati iteration coverges to P^* starting everywhere in the inerval (\tilde{P}, ∞) as shown in the figure. \square

Random System Matrices Suppose that $\{A_0, B_0\}, \dots, \{A_{N-1}, B_{N-1}\}$ are not known but *independent random matrices* that are also independent of the ω_k . The DP algorithm is given by:

$$J_N(x_N) = x_N^T Q_N x_N$$

$$J_k(x_k) = \min_{u_k} \mathbb{E}_{\omega_k, A_k, B_k} \left\{ x_k^T Q_k x_k + u_k^T R_k u_k + J_{k+1}(A_k x_k + B_k u_k + \omega_k) \right\}$$

Following the similar proof in Theorem (3.1), we imply

$$\mu^* k(x_k) = L_k x_k, \text{ where } L_k = -(R_k + \mathbb{E}\{B_k^T K_{k+1} B_k\})^{-1} \mathbb{E}\{B_k^T K_{k+1} A_k\}$$

and the cost-to-go function $J_k(x_k) = x_k^T K_k x_k + \text{constant}$, with

$$K_N = Q_N$$

$$K_k = \mathbb{E}\{A_k^T K_{k+1} A_k\} - \mathbb{E}\{A_k^T K_{k+1} B_k\} \\ (R_k + \mathbb{E}\{B_k^T K_{k+1} B_k\})^{-1} \mathbb{E}\{B_k^T K_{k+1} A_k\} + Q_k$$

However, there are several drawbacks for random matrices systems:

1. Certainty Equivalence may not hold (what is certainty equivalence?) Formally speaking, the optimal controller may not be the same as for the corresponding deterministic problem, when the randomness is replaced by its expectation.
2. Riccati equation may not converge to a steady-state.

Consider the scalar stationary system such that for the update for K_{k+1} , the random matrices A_k, B_k are replaced by their expectation. Therefore, the update for K_{k+1} can be written as

$$\begin{cases} K_{k+1} = \tilde{F}(K_k) \\ \tilde{F}(K) = \frac{\mathbb{E}\{A^2\}BK}{\mathbb{E}\{B^2\}K + R} + Q + \frac{TK^2}{\mathbb{E}\{B^2\}K + R} \\ T = \mathbb{E}\{A^2\}\mathbb{E}\{B^2\} - (\mathbb{E}\{A\})^2(\mathbb{E}\{B\})^2 \end{cases}$$

When $T = 0$, the problem reduces to the case where A, B are not random; and we can obtain the well-studied asymptotic behavior; the convergece also happens as T is small. However, for large enough T , as shown in Fig (3.2), the graph of F do not intersect with the 45-degree line on the positive interval, i.e., the Riccati equation diverges to infinity.




Figure 3.2: Graphical Illustration for a scalar stationary system

3.2 Inventory Control

- State: inventory level x_k
- Control: number of orders u_k
- Disturbance: demand ω_k
- System dynamics: $x_{k+1} = x_k + u_k - \omega_k$, $k = 0, 1, \dots, N - 1$
- Stage cost:

$$c \cdot u_k + H(x_k + u_k)$$

where

$$H(x_k + u_k) := \mathbb{E}_{\omega_k} \left[p \cdot (x_k + u_k - \omega_k)^- + h \cdot (x_k + u_k - \omega_k)^+ \right]$$

The goal is to minimize the total cost (question: do we need to add the expectation symbol?)

$$\mathbb{E} \left\{ \sum_{k=0}^{N-1} (cu_k + H(x_k + u_k)) \right\}$$

The dynamic programming algorithm applies as follows:

$$J_N(x_N) = 0$$

$$J_k(x_k) = \min_{u_k \geq 0} [c \cdot u_k + H(x_k + u_k) + \mathbb{E}_{\omega_k} J_{k+1}(x_{k+1})]$$

Optimal Policy Regarding $y_k = x_k + u_k$ as the decision variable, it suffices to solve

$$\min_{y_k \geq x_k} G_k(y_k) - cx_k,$$

where $G_k(y) = cy + H(y) + \mathbb{E}_{\omega} [J_{k+1}(y - \omega)]$

Proposition 3.1. If G_k is convex, and $\lim_{|x| \rightarrow \infty} G_k(x) \rightarrow \infty$, we have the optimal policy

$$\mu_k^*(x_k) = \begin{cases} S_k - x_k, & \text{if } x_k < S_k \\ 0, & \text{if } x_k \geq S_k \end{cases}$$

where S_k is the minimizer of $G_k(y)$.

Now we show the assumption for proposition (3.1) mostly holds. We need to add the assumption $c < p$, since otherwise it would never be optimal to order new items at the last period.

Proposition 3.2. Assume that $c < p$, then J_k is convex for all k and

$$\lim_{|x| \rightarrow \infty} J_k(x) \rightarrow \infty, \quad k = 0, 1, \dots, N-1$$

Proof. 1. It's easy to see that $J_N \equiv 0$ is convex.

2. Since $c < p$ and $H'(y) \rightarrow -p$ as $y \rightarrow -\infty$, we imply that $G_{N-1} := cy + H(y)$ has a derivative that becomes negative as $y \rightarrow -\infty$ and becomes positive when $y \rightarrow \infty$ (see Fig (3.3)). Therefore,

$$\lim_{|y| \rightarrow \infty} G_{N-1}(y) = \infty.$$

As have shoen,

$$\mu_{N-1}^*(x_{N-1}) = \begin{cases} S_{N-1} - x_{N-1}, & x_{N-1} < S_{N-1} \\ 0, & x_{N-1} \geq S_{N-1} \end{cases}$$




Figure 3.3: Graphical Illustration for the structure of cost-to-go function

Or equivalently, $\mu_{N-1}^*(x_{N-1}) = (S_{N-1} - x_{N-1})^+$ (question: can we write in this form?). Substituting μ_{N-1}^* into J_{N-1} , we imply

$$J_{N-1}(x_{N-1}) = \begin{cases} c \cdot (S_{N-1} - x_{N-1}) + H(S_{N-1}), & \text{if } x_{N-1} < S_{N-1} \\ H(x_{N-1}), & \text{if } x_{N-1} \geq S_{N-1} \end{cases}$$

which is clearly a convex function (question: why?)

Therefore, given the convexity of J_N , we are able to show the convexity of J_{N-1} . Furthermore,

$$\lim_{|y| \rightarrow \infty} J_{N-1}(y) = \infty.$$

3. Similarly, given that J_{k+1} is convex and $\lim_{|y| \rightarrow \infty} J_{k+1}(y) = \infty$, $\lim_{|y| \rightarrow \infty} G_k(y) = \infty$, for $k = N-2, \dots, 0$, we imply

$$J_k(x_k) = \begin{cases} c \cdot (S_k - x_k) + H(S_k) + \mathbb{E}\{J_{k+1}(S_k - \omega_k)\}, & \text{if } x_k < S_k \\ H(x_k) + \mathbb{E}\{J_{k+1}(x_k - \omega_k)\}, & \text{if } x_k \geq S_k \end{cases}$$

where S_k is the unconstrained minimizer of G_k . Moreover, we imply J_k is convex, $\lim_{|y| \rightarrow \infty} J_k(y) = \infty$ and $\lim_{|y| \rightarrow \infty} G_{k-1}(y) = \infty$. \square

Therefore, we reduce the policy-based optimization into the pattern of finding the baseline S_k , i.e., a traditional parametric optimization problem.

Proposition 3.3. If the disturbance ω_k is i.i.d., then we imply the stationary optimal policy, i.e., the policy baseline S_k^* becomes constant.

Now suppose further that

$$c(u_k) = \begin{cases} K + c \cdot u_k, & u_k > 0 \\ 0, & u_k = 0 \end{cases}$$

Now the question is that do we obtain the pattern with baseline?

$$\begin{aligned} J_k(x_k) &= \min_{u_k > 0} [K + cu_k + H(x_k + u_k)] + \mathbb{E}_{\omega_k} J_{k+1}(x_{k+1}) \\ &\quad , 0 + H(x_k) + \mathbb{E} J_{k+1}(x_k)] \\ &= \min \left[\min_{y_k \geq x_j} (K + G_k(y_k)), G_k(x_k) \right] - cx_k \end{aligned}$$

where $f(x_k) := \min_{y_k \geq x_j} (K + G_k(y_k))$ has its shape. Then we need to compare the two curves $f(x_k)$ and $G_k(x_k)$:

$$u_k^* = \begin{cases} S_k - x_k, & \text{if } x_k \leq S_k \\ 0, & \text{if } x_k > S_k \end{cases}$$

- If the inventory is very low (compare the baseline s_k), order $u_k^* = S_k - x_k$ to make the inventory into S_k
- If not that low, not order anything.

Remark 3.1. Unfortunately, $J_{k+1}(x_k)$ is convex does not necessarily imply $J_k(x_k)$ is convex. There is a concept K -convexity in literature. Under this K -convexity setting, we can show that the control u_k^* still obtain the similar pattern but with two baselines, i.e., the decision can be made by first computing the baseline (s, S) , and then obtaining the (s, S) strategy.

The key for reducing computational complexity is to find patterns for decision variable u_k^* .

3.3 Optimal Stopping Problem

There are two possible controls for a stopping problem:

1. Stop, i.e., incur a one-time stopping cost, and move to a *cost-free, absorbing* stop state.
2. Continue, using $x_{k+1} = f_k(x_k, \omega_k)$ and incurring the *cost-per-stage*

Therefore, eacy policy should consist of a partition of a set of states x_k into two regions:

1. Stop region: where we stop?
2. Continue region, where do we continue?

The natural question is that can we obtain an optimal policy that is similar to the pattern of the *baseline* studied in inventory control problem?

Asset Selling Suppose an industry wants to sell its items. This industry receives a random price ω_k for $k = 1, \dots, N - 1$, which are assumed to be i.i.d. with known distribution. This industry may accept ω_k and invest the money at a fixed rate of interest r ; or reject ω_k and wait for ω_{k+1} . It must accept the last offer ω_{N-1}

We may formulate this problem as follows:

- State: $x_k = \omega_k$
- Decision variable: u^1 for sell and u^2 for wait.
- Stage reward:

$$g(x_k, u_k, \omega_k) = \begin{cases} x_k(1+r)^{N-1}, & \text{if } u_k = u^1 \\ 0, & \text{if } u_k = u^2 \end{cases}$$

- Therefore, the DP algorithm is by solving the optimization problem below recursively:

$$J_N(x_N) = \begin{cases} x_N & \text{if } x_N \neq T \\ 0, & \text{if } x_N = T \end{cases}$$

$$J_k(x_k) = \begin{cases} \max \left[(1+r)^{N-k} x_k, \mathbb{E}\{J_{k+1}(\omega_k)\} \right] & \text{if } x_k \neq T \\ 0, & \text{if } x_k = T \end{cases}$$

Specifically, note that for k -th stage, if $x_k(1+r)^{N-k} \geq \mathbb{E}_{\omega_k}[J_{k+1}(\omega_k)]$, then the industry should decide to sell; otherwise decide to wait.

Define the variable (baseline)

$$\alpha_k = \frac{\mathbb{E}J_{k+1}(\omega)}{(1+r)^{N-k}},$$

then the optimal policy should be:

$$\begin{aligned} &\text{accept the offer } x_k && \text{if } x_k > \alpha_k \\ &\text{reject the offer } x_k && \text{if } x_k < \alpha_k \end{aligned}$$

In particular, the high baseline corresponds to the high standard for offering the price.

Future Analysis

Theorem 3.3. Based on the assumption of i.i.d. ω_k , the baseline $\alpha_k \geq \alpha_{k+1}$ for all k .

Proof. First introduce the functions

$$V_k(x_k) = \frac{J_k(x_k)}{(1+r)^{N-k}}, \quad x_k \neq T$$

The DP algorithm reduces into

$$\begin{aligned} V_N(x_N) &= x_N \\ V_k(x_k) &= \max \left[x_k, (1+r)^{-1} \mathbb{E}_{\omega} \{ V_{k+1}(\omega) \} \right] \end{aligned}$$

Note that $\alpha_k = \mathbb{E}\{V_{k+1}(\omega)\}/(1+r)$. It's clear that

$$V_{N-1}(x) \geq V_N(x), \quad \forall x \geq 0$$

Following the similar procedure, we derive $V_k(x) \geq V_{k+1}(x)$ for all x and k . Since $\alpha_k = \mathbb{E}\{V_{k+1}(\omega)\}/(1+r)$, we obtain $\alpha_k \geq \alpha_{k+1}$, as desired. \square

Remark 3.2. We can also show that $\alpha_k \rightarrow \bar{\alpha}$ as $k \rightarrow \infty$, where $\bar{\alpha}$ is a constant satisfying

$$(1+r)\bar{\alpha} = P(\bar{\alpha})\bar{\alpha} + \int_{\bar{\alpha}}^{\infty} \omega \, dP(\omega)$$

Therefore, for an infinite horizon, the optimal policy is stationary.

Now we show that $\{\alpha_k\}$ is monotonically decreasing.

Define

$$V_k(x_k) = \frac{J_k(x_k)}{(1+r)^{N-k}} \implies V_k(x_k) = \max\left(x_k, \frac{\mathbb{E}V_{k+1}(\omega)}{1+r}\right)$$

and $V_N(x_N) = x_N$.

Note that we show $V_{k+1}(x) \leq V_k(x)$ first:

$$\begin{aligned} V_N(x) &= x \\ V_{N-1}(x) &= \max(x, \mathbb{E}V_N(\omega)) \end{aligned}$$

It's clear that $V_N(x) \leq V_{N-1}(x)$.

$$\begin{aligned} V_K(x) &= \max\left(x, \frac{\mathbb{E}V_{k+1}(\omega)}{1+r}\right) \\ V_{k-1}(x) &= \max\left(x, \frac{\mathbb{E}V_k(\omega)}{1+r}\right) \end{aligned}$$

It's clear that $V_k(x) \leq V_{k-1}(x)$ for any k . Therefore, $\{\alpha_k\}$ is monotonically decreasing.

Note that the key for this proof is the i.i.d. of ω_k .

- Pattern
- Trends
- Convergence: If there are infinite horizons, what is the bottleneck for the policy.

In this case,

$$\begin{aligned} \alpha_k &= \frac{\mathbb{E}V_{k+1}(\omega)}{1+r} = \frac{1}{1+r} \left[\int_0^{\alpha_{k+1}} \alpha_{k+1} dP(\omega) + \int_{\alpha_{k+1}}^{\infty} \omega dP(\omega) \right] \\ &= \frac{P(\alpha_{k+1})}{1+r} \alpha_{k+1} + \frac{1}{1+r} \int_{\alpha_{k+1}}^{\infty} \omega dP(\omega) \end{aligned} \tag{3.1}$$

where the $P(\alpha_{k+1})$ is the cdf, and α_k is bounded.

Start from $\alpha_N = 0$, derive $\alpha_{N-1}, \dots, \alpha_0$ sequentially.

When $N \rightarrow \infty$, note that computing the equation above is based on the derivation of α_{N+1} . We can imply the convergence of α , and we can compute α_{N+1} based on (3.1).

4

Dynamic Programming with imperfect information

4.1 Problem Setting

- State: x_k
- Control: $u_k \in \mu_k(x_k)$
- Disturbance: ω_k
- Dynamics: $x_{k+1} = f_k(x_k, u_k, \omega_k)$
- stage cost: $g_N(x_N), g_k(x_k, u_k, \omega_k)$

Imperfect Information Setting

- Observation: Z_k
- measurement: $z_k = h_k(x_k, u_{k-1}, v_k)$, where v_k denotes the measurement noise, and v_k is dependent on $x_0, \dots, x_{k-1}, u_0, \dots, u_{k-1}, v_0, \dots, v_{k-1}$
- Information vector: $I_k = (z_0, \dots, z_k, u_0, \dots, u_{N-1})$
- The goal is to derive the optimal policy

$$\pi = (\mu_0(I_0), \dots, \mu_{N-1}(I_{N-1}))$$

where $u_k^* = \mu_k(I_k)$ to minimize the cost function

$$J_\pi = \mathbb{E}_{\omega_0, \dots, \omega_{N-1}, v_0, \dots, v_{N-1}, x_0} \left\{ \sum_{k=0}^{N-1} g_k(x_k, \mu_k(I_k), \omega_k) + g_N(x_N) \right\}$$

Therefore, we need the distribution of x_0 .

4.2 Reformulation as a perfect state information system

- State: $I_k = (z_0, \dots, z_k, u_0, \dots, u_{k-1})$
- Control: $u_k \in U_k$
- Disturbance: z_{k+1}
- System dynamics:

$$I_{k+1} = (I_k, u_{k+1}, z_{k+1})$$

where z_{k+1} is the unknown term, considered as the noise

- Stage cost:

$$\mathbb{E}_{x_k, \omega_k} (g_k(x_k, u_k, \omega_k) \mid I_k, u_k) = \sum_{x_k} \mathbb{E}_{\omega_k} g_k(x_k, u_k, \omega_k) \cdot P(x_k \mid I_k)$$

Therefore, this stage cost is a function w.r.t. I_k, u_k .

Dynamic algorithm applies as follows:

$$\begin{aligned} J_{N-1}(I_{N-1}) &= \min_{u_{N-1} \in \mu_{N-1}} \mathbb{E}_{x_{N-1}, \omega_{N-1}} [g_{N-1}(x_N, u_{N-1}, \omega_{N-1}) \mid I_{N-1}, \omega_{N-1}] \\ &\quad + \mathbb{E}_{\omega_{N-1}, x_{N-1}} [g_N(x_N) \mid I_{N-1}, u_{N-1}] \end{aligned}$$

For $k = N-1, N-2, \dots, 1$,

$$\begin{aligned} J_k(I_k) &= \min_{u_k \in U_k} \mathbb{E}_{x_k, \omega_k} [g_k(x_k, u_k, \omega_k) \mid I_k, u_k] \\ &\quad + \mathbb{E}_{z_{k+1}} [J_{k+1}(I_k, u_{k+1}, z_{k+1}) \mid I_k, u_k] \end{aligned}$$

We need the conditional distribution $P(x_k \mid I_k)$ and $P(z_{k+1} \mid I_k)$.

4.3 Machine Repair Example

- State: x_k is P or \bar{P}
- Observation: $Z_k = \{G, B\}$
- Control: $u_k = \{C, S\}$, where C denotes continue, S denotes stop to check.

- Disturbance:

System Dynamics:

$$P(x_{k+1} = (P, \bar{P}) \mid x_k = P) = (2/3, 1/3)$$

$$P(x_{k+1} = (P, \bar{P}) \mid x_k = \bar{P}) = (0, 1)$$

- Stage cost:

$$g(x_k, u_k) = \begin{cases} 0, & (x_k, u_k) = (P, C) \\ 2, & (x_k, u_k) = (\bar{P}, C) \\ 1, & (x_k, u_k) = (P, S) \\ 1, & (x_k, u_k) = (\bar{P}, S) \end{cases}$$

- Measurement:

$$P(Z_k = G \mid x_k = P) = 3/4, \quad P(Z_k = B \mid x_k = \bar{P}) = 3/4$$

- Cost-to-go Function:

$$N = 2, \quad J_N(I_N) = 0$$

$$\begin{aligned} J_1(I_1) &= \min_{u_1 \in \{C, S\}} \mathbb{E}_{x_1} [g(x_1, u_1) \mid I_1, u_1] + \underbrace{\mathbb{E} J_2}_0 \\ &= \min_{u_1} \{ \mathbb{E}_{x_1} [g(x_1, c) \mid I_1, c], \mathbb{E}_{x_1} [g(x, s) \mid I_1, s] \} \end{aligned}$$

We need the conditional distribution $P(x_k \mid I_k)$, and the initial distribution $P(x_0 = P) = 2/3, P(x_0 = \bar{P}) = 1/3$.

Prior: $P(x_0)$; conditional distribution: $P(z_0 \mid x_0)$; posterior distribution: $P(x_0 \mid z_0)$. Marginal: $P(z_0)$

Bayes's Theorem:

$$P(x_0 \mid z_0) = \frac{P(x_0, z_0)}{P(z_0)} = \frac{P(z_0 \mid x_0)P(x_0)}{\sum_{x_0} P(z_0 \mid x_0)P(x_0)}$$

Therefore,

$$P(x_0 = P \mid z_0 = G) = \frac{P(x_0 = P, z_0 = G)}{p(z_0 = G)} = \frac{2/3 \times 3/4}{2/3 \times 3/4 + 1/3 \times 1/4} = 6/7$$

$$P(x_1 \mid I_1) = P(x_1 \mid z_0, z_1, u_1) = \frac{P(z_1 \mid x_1, z_0, u_1) \cdot P(x_1 \mid z_0, u_0)}{P(z_1 \mid z_0, u_1)}$$

5

The Distribution and Installation

5.1 Pre-requisites

You will need a working LaTeX installation. We recomend using pdfflatex to process the files. You will also need biber.exe installed. This is distributed as part of the latest versions of LiveTex and MikTex. If you have problems, please let us know.

5.2 The Distribution

The distribution contains 2 folders: `nowfnt` and `nowfnttexmf`.

5.2.1 Folder `nowfnt`

This folder contains the following files using a flat stucture required to compile a FnT issue:

- `essence_logo.eps`
- `essence_logo.pdf`
- `now_logo.eps`
- `now_logo.pdf`

- nowfnt.cls
- nowfnt-biblatex.sty
- NOWFnT-data.tex

It also contains the following folders:

journaldata A set of data files containing the journal-specific data for each journal. There are three files per journal:

<jrnlcode>-editorialboard.tex

<jrnlcode>-journaldata.tex

<jrnlcode>-seriespage.tex

<jrnlcode> is the code given in Appendix A. You will need these three files to compile your article.

SampleArticle This folder contains this document as an example of an article typeset in our class file. The document is called `FnTarticle.tex`. It also contains this PDF file and the `.bib` file.

5.2.2 Folder `nowfnttexmf`

This folder contains all the files required in a `texmf` structure for easy installation.

5.3 Installation

If your \LaTeX installation uses a `localtexmf` folder, you can copy the `nowtexmf` folder to the `localtexmf` folder and make it known to your \TeX installation. You can now proceed to use the class file as normal.

If you prefer to use the flat files, you will need to copy all the required files each time into the folder in which you are compiling the article. Do not forget to copy the three data files for the specific journal from the folder `journaldata`.

You may need to configure your \TeX editor to be able to run the programs. If you have problems installing these files in your own system, please contact us. We use Computer Modern fonts for some of the

journals. You will need to make sure that these fonts are installed. Refer to your system documentation on how to do this.

6

Quick Start

The now-journal class file is designed in such a way that you should be able to use any commands you normally would. However, do **not** modify any class or style files included in our distribution. If you do so, we will reject your files.

The preamble contains a number of commands for use when making the final versions of your manuscript once it has been accepted and you have been instructed by our production team.

6.1 `\documentclass`

The options to this command enable you to choose the journal for which you producing content and to indicate the use of biber.

```
\documentclass[<jrnlcode>,biber]{nowfnt}.
```

`<jrnlcode>` is the pre-defined code identifying each journal. See Appendix A for the appropriate `<jrnlcode>`.

6.2 `\issuesetup`

These commands are only used in the final published version. Leave these as the default until our production team instructs you to change them.

6.3 `\maintitleauthorlist`

This is the authors list for the cover page. Use the name, affiliation and email address. Separate each line in the address by `\\`.

Separate authors by `\and`. Do not use verbatim or problematic symbols. `_` (underscore) in email address should be entered as `_`. Pay attention to long email addresses.

If your author list is too long to fit on a single page you can use double column. In this case, precede the `\maintitleauthorlist` command with the following:

```
\booltrue{authortwocolumn}
```

6.4 `\author` and `\affil`

These commands are used to typeset the authors and the affiliations on the abstract page of the article and in the bibliographic data.

`\author` uses an optional number to match the author with the affiliation. The author name is written `<surname>`, `<firstname>`.

`\affil` uses an optional number to match the author with the author name. The content is `<affiliation>`; `<email address>`.

6.5 `\addbibresource`

Use this to identify the name of the bib file to be used.

7

Style Guidelines and L^AT_EX Conventions

In this section, we outline guidelines for typesetting and using L^AT_EX that you should follow when preparing your document

7.1 Abstract

Ensure that the abstract is contained within the

`\begin{abstract}`

environment.

7.2 Acknowledgements

Ensure that the acknowledgements are contained within the

`\begin{acknowledgements}`

environment.

7.3 References

now publishers uses two main reference styles. One is numeric and the other is author/year. The style for this is pre-defined in the L^AT_EX

distribution and must not be altered. The style used for each journal is given in the table in Appendix A. Consult the sample-now.bib file for an example of different reference types.

The References section is generated by placing the following commands at the end of the file.

```
\backmatter
\printbibliography
```

7.4 Citations

Use standard `\cite`, `\citep` and `\citet` commands to generate citations.

Run biber on your file after compiling the article. This will automatically create the correct style and format for the References.

7.4.1 Example citations

This section cites some sample references for your convenience. These are in author/year format and the output is shown in the References at the end of this document.

Example output when using `citet`: **arvolumenumber** is a citation of reference 1 and Bertsekas (1995) is a citation of reference 2.

Example output when using `citep`: (**beditorvolumenumber**) is a citation of reference 3 and (**inproceedings**) is a citation of reference 4.

7.5 Preface and Other Special Chapters

If you want to include a preface, it should be defined as follows:

```
\chapter*{Preface}
\markboth{\sffamily\slshape Preface}
{\sffamily\slshape Preface}
```

This ensures that the preface appears correctly in the running headings.

You can follow a similar procedure if you want to include additional unnumbered chapters (*e.g.*, a chapter on notation used in the paper), though all such chapters should precede Chapter 1.

Unnumbered chapters should not include numbered sections. If you want to break your preface into sections, use the starred versions of `section`, `subsection`, *etc.*

7.6 Long Chapter and Section Names

If you have a very long chapter or section name, it may not appear nicely in the table of contents, running heading, document body, or some subset of these. It is possible to have different text appear in all three places if needed using the following code:

```
\chapter[Table of Contents Name]{Body Text Name}  
\chaptermark{Running Heading Name}
```

Sections can be handled similarly using the `sectionmark` command instead of `chaptermark`.

For example, the full name should always appear in the table of contents, but may need a manual line break to look good. For the running heading, an abbreviated version of the title should be provided. The appearance of the long title in the body may look fine with L^AT_EX's default line breaking method or may need a manual line break somewhere, possibly in a different place from the contents listing.

Long titles for the article itself should be left as is, with no manual line breaks introduced. The article title is used automatically in a number of different places by the class file and manual line breaks will interfere with the output. If you have questions about how the title appears in the front matter, please contact us.

7.7 Internet Addresses

The class file includes the `url` package, so you should wrap email and web addresses with `\url{}`. This will also make these links clickable in the PDF.

8

Compiling Your FnT Article

During the first run using the class file, a number of new files will be created that are used to create the book and ebook versions during the final production stage. You can ignore these until preparing the final versions as described in Section 8.3. A complete list of the files produced are given in Appendix B.

8.1 Compiling Your Article Prior to Submission

To compile an article prior to submission proceed as follows:

Step 1: Compile the L^AT_EX file using pdf_latex.

Step 2: Run biber on your file.

Step 3: Compile again using pdf_LaT_EX. Repeat this step.

Step 4: Inspect the PDF for bad typesetting. The output PDF should be similar to FnTarticle.pdf. Work from the first page when making adjustments to resolve bad line breaks and bad page breaks. Re-run pdf_LaT_EX on the file to check the output after each change.

8.2 Preparing the Final Versions

If you choose the option to compile the final versions of your PDF for publication, you will receive a set of data from our production team upon final acceptance. With the exception of "lastpage", enter the data into the `\issuesetup` command in the preamble.

lastpage= This is the last page number in the sequential numbering of the journal volume. You will need to enter this once you have compiled the article once (see below).

8.3 Compiling The Final Versions

The final versions should be created once you received all the bibliographic data from our Production Team and you've entered it into the preamble. You will be creating a final online journal version pdf; a printed book version pdf; and an ebook version pdf.

Step 1: Compile the L^AT_EX file using pdfLaTeX.ArAuthor *et al.*, 2014

Step 2: Run biber on your file.

Step 3: Compile again using pdfLaTeX. Repeat this step.

Step 4: Inspect the PDF for bad typesetting. The output PDF should be similar to FnTarticle.pdf. Work from the first page when making adjustments to resolve bad line breaks and bad page breaks. Re-run pdfLaTeX on the file to check the output after each change.

Step 5: When you are happy with the output make a note of the last page number and enter this in `\issuesetup`.

Step 6: Compile the article again.

Step 7: Open the file `<YourFilename>-nowbook.tex`. This will generate the printed book version pdf.

Step 8: Compile the L^AT_EX file using pdfflatex.

Step 9: Run biber on your file.

Step 10: Compile again using pdfLaTeX. Repeat this step.

Step 11: Repeat steps 7-10 on the file: <YourFilename>-nowebook.tex.
This will generate the ebook version pdf.

Step 12: Repeat steps 7-10 on the file: <YourFilename>-nowplain.tex.
This will generate a plain version pdf of your article. If you intend to post your article in an online repository, please use this version.

8.4 Wednesday

Proposition 8.1. G is hamiltonian implies that for any nonempty $S \subseteq V$, $G - S$ has at most $|S|$ components.

Theorem 8.1. G is Hamiltonian if for any vetices v, w such that $(v, w) \notin E$,

$$\deg(v) + \deg(w) \geq n$$

Knapsacle Problem

$$\begin{array}{ll} \max & \sum_{i=1}^n v_i u_i \\ \text{such that} & \sum_{i=1}^n w_i u_i \leq W \\ & u_i \in \{0, 1\}, \quad i = 1, \dots, n \end{array}$$

This is the binary integer programming problem, which is NP-complete. We cannot find exact solution to this problem. There are 2^n possible solutions, which requires exponential time.

We can find some “pseudo”-polynomial algorithm. Assumption:

1. W is an integer.

State: remaining capacity, define as X_k .

Stage: item $1, \dots, n$.

$$J_N(x_N) = \begin{cases} v_N & \text{if } N \leq X_N \\ 0, & \text{if } w_N > X_N \end{cases}$$

$$J_N(x_{N-1}) = \begin{cases} 0 + J_N(X_N), & \text{if } w_{N-1} > X_{N-1}, \text{ Here } X_N = \\ \max\{v_{N-1} + J_N(x_N), 0 + J_N(x_N)\}, & \text{if } w_{N-1} \leq X_{N-1} \end{cases}$$

Here why the DP has “pesduo”-polynomial time, and the w_i should be integer?

Let's list states $\{x_1, \dots, x_N\}$. For each stage k , there would be $W+1$.

Each iteration at most 2 computation, and therefore we face $2(W+1) \cdot N$

8.4.1 Label Correcting Methods

Shortest Path. Back up some information.

The label refers to the intermediate information.

A^* -algorithm; Bellman-Ford Algorithm

Benchmark: MILP, CPLEX, CVX.

8.4.2 DP problems with perfect state information

Linear Quadratic System Let $x_k \in \mathbb{R}^1$ be the state; $u_k \in \mathbb{R}^n$ be the control; $\omega_k \in \mathbb{R}^n$ be the disturbance.

System Dynamics.

$$x_{k+1} = A_k x_k + B_k u_k + \omega_k$$

Stage cost:

$$x'_k Q_k x_k + u'_k R_k u_k$$

Here Q_k is required to be positive-semi-definite symmetric matrix; R_k to be positive-definite symmetric matrix.

$$J_{n-1}(x_{n-1}) = \min_{u_{n-1}} \mathbb{E}_\omega \{x'_{n-1} Q_{n-1} x_{n-1} + u'_{n-1} R_{n-1} u_{n-1} + J_n(x_n)\}$$

where

$$J_n(x_n) = (A_{n-1} x_{n-1} + B u_{n-1} + \omega)' Q_n (A_{n-1} x_{n-1} + B u_{n-1} + \omega)$$

8.4.3 Linear Quadratic

Dynamics:

$$x_{k+1} = A_k x_k + B_k u_k + \omega_k$$

Stage cost:

$$x'_k Q_k x_k + u'_k R_k u_k, \quad Q_k \succeq 0, R_k \succ 0$$

Cost to go function:

$$J_k(x_k) = \min_{u_k} \mathbb{E}_{\omega_k} [x'_k Q_k x_k + u'_k R_k u_k + J_{k+1}(x_{k+1})]$$

The final cost is

$$J_N(x_N) = x'_N Q_N x_N$$

At $N - 1$ stage

$$\begin{aligned} & x'_{N-1} Q_{N-1} x_N + u'_{N-1} R_{N-1} u_{N-1} \\ & + \mathbb{E}[(A_{N-1} x_{N-1} + B_{N-1} u_{N-1} + \omega_{N-1})' Q_N (A_{N-1} x_{N-1} + B_{N-1} u_{N-1} + \omega_{N-1})] \end{aligned}$$

Optimization model:

$$\frac{\partial}{\partial u} (u' H u + 2r' u + c) = 2H u + 2r \implies H u = -r \implies u^* = -H^{-1} r$$

For $N - 1$ stage, we have

$$\begin{aligned} H &= R_{N-1} + B'_{N-1} Q_N B_{N-1} \\ r' &= x'_{N-1} A'_{N-1} Q_N B_N \\ c &= x'_{N-1} (Q_{N-1} + A'_{N-1} Q_N A_{N-1}) x_{N-1} + \mathbb{E}_{\omega} (\omega'_{N-1} Q_N \omega_{N-1}) \end{aligned}$$

Therefore,

$$u_{N-1}^* = -(R_{N-1} + B'_{N-1} Q_N B_{N-1})^{-1} B_N Q_N A_{N-1} x_{N-1},$$

which is linear in terms of x_{N-1} , which is so called linear controller.

Therefore,

$$J_{N-1}(x_{N-1}) = x'_{N-1} K_{N-1} x_{N-1} + \mathbb{E}(\omega'_{N-1} Q_N \omega_{N-1})$$

The linear controller will be tested during mid-term or final.

$$J_0(x_0) = x'_0 K_0 x_0 + \sum_{k=0}^{N-1} \mathbb{E}_{\omega} \{ \omega'_k K_{k+1} \omega_k \}$$

$$K_{k-1} = f(K_k) \implies K_{k-1} = K_k, \text{ for large } k.$$

Riccati Equation for stationary system, i.e., $A_k = A, B_k = B, Q_k = Q, R_k = R$.

$$K = A'(K - KB(B'KB + R)^{-1}B'K)A + Q$$

Therefore as time goes long enough, the cost to go function $J_k(x_k)$ will be a constant plus over stages. such a constant K is called the *optimal stationary controller*.

Stability. For $u^* = Lx$, we imply

$$x_{k+1} = Ax_k + Bu_k + \omega_k = (A + BL)x_k + \omega_k$$

Care about $\lim_{k \rightarrow \infty} (A + BL)^k = 0$.

Here

$$L = -(B'KB + R)^{-1}B'KA$$

It suffices to solve

$$P_{k+1} = A^2 \left(P_k - \frac{B^2 P_k^2}{B^2 P_k + R} \right) + Q$$

Then consider the case that A_k, B_k are all random matrices, i.e., independent. $Q_k \succeq 0, R_k \succ 0$.

$$L_k = -(R_k + \mathbb{E}(B'_k K_{k+1} B_k))^{-1} \mathbb{E}(B'_k K_{k+1} A_k)$$

Note that

$$P_\infty = \frac{\mathbb{E}A^2 R P}{R + \mathbb{E}B^2 P} + \frac{\mathbb{E}A^2 \mathbb{E}B^2 - (\mathbb{E}A)^2 (\mathbb{E}B)^2}{R + \mathbb{E}B^2 P}$$

Certainty Equivalence

State: x_k , inventory level

Control: $u_k \geq 0$, number of orders placed

Disturbance: ω_k : demand

Dynamics:

$$x_{k+1} = (x_k + u_k - \omega_k)$$

Stage cost:

- Ordering cost: $c \cdot u_k$

- Maintaining cost: full backlog. $\mathbb{E}_{\omega_k}(r(x_k + u_k - \omega_k))$.

where $r(z) = hz^+ + pz^-$ is a convex function.

We imply the maintaining cost is $H(x_k + u_k)$, which is convex.

$$J_k(x_k) = \min_{u_k \geq 0} \{cu_k + H(x_k + u_k) + \mathbb{E}[J_{k+1}(x_k + u_k - \omega_k)]\}$$

Define $Y = x_k + u_k$, and therefore

$$G(Y) = cY + H(Y) + \mathbb{E}[J(Y - \omega_k)]$$

Therefore, we can always set $Y = x_k + u_k = S_k$, where S_k minimizes $G(Y)$.

Acknowledgements

The authors are grateful to Ulrike Fischer, who designed the style files, and Neal Parikh, who laid the groundwork for these style files.

Appendices

A

Journal Codes

The table below shows the journal codes to be used in `\documentclass`.
 For Example: `\documentclass[ACC,biber]{nowfnt}`

Journal	<jrnocode>	Ref. Style
Annals of Corporate Governance	ACG	Author/Year
Annals of Science and Technology Policy	ASTP	Author/Year
FnT Accounting	ACC	Author/Year
FnT Comm. and Information Theory	CIT	Numeric
FnT Databases	DBS	Author/Year
FnT Econometrics	ECO	Author/Year
FnT Electronic Design Automation	EDA	Author/Year
FnT Electric Energy Systems	EES	Author/Year
FnT Entrepreneurship	ENT	Author/Year
FnT Finance	FIN	Author/Year
FnT Human-Computer Interaction	HCI	Author/Year
FnT Information Retrieval	INR	Author/Year
FnT Information Systems	ISY	Author/Year
FnT Machine Learning	MAL	Author/Year
FnT Management	MGT	Author/Year
FnT Marketing	MKT	Author/Year
FnT Networking	NET	Numeric
		<i>Continues</i>

Journal	<jrnocode>	Ref. Style
FnT Optimization	OPT	Numeric
FnT Programming Languages	PGL	Author/Year
FnT Robotics	ROB	Author/Year
FnT Privacy and Security	SEC	Author/Year
FnT Signal Processing	SIG	Numeric
FnT Systems and Control	SYS	Author/Year
FnT Theoretical Computer Science	TCS	Numeric
FnT Technology, Information and OM	TOM	Author/Year
FnT Web Science	WEB	Author/Year

B

Files Produced During Compilation

The files that are created during compilation are listed below. The additional *.tex files are used during the final production process only. See Section 8.3.

```
<YourFilename>-nowbook.tex  
<YourFilename>-nowchapter.tex  
<YourFilename>-nowebook.tex  
<YourFilename>-nowechapter.tex  
<YourFilename>-nowsample.tex  
<YourFilename>-nowplain.tex  
<YourFilename>.aux  
<YourFilename>.bbl  
<YourFilename>.bcf  
<YourFilename>.blg  
<YourFilename>.log  
<YourFilename>.out  
<YourFilename>.pdf  
<YourFilename>.run.xml  
<YourFilename>.synctex.gz  
<YourFilename>.tex  
<YourFilename>.toc
```

References

- ArAuthor, A., B. Author, and C. Author (2014). “An article on something interesting (volume only)”. *Journal of interesting Things*. 6: 1–122. I have something special to say about this publication. ISSN: 0899-8256. DOI: 10.5161/202.00000013. URL: <http://www.nowpublishers.com/> (accessed on 07/10/2014).
- Bertsekas, D. (1995). “Dynamic Programming and Optimal Control”. In: vol. 1.
- Sutton, R. S. (1988). “Learning to predict by the methods of temporal differences”. *Machine Learning*. 3: 9–44.
- Watkins, C. J. C. H. and P. Dayan (1992). “Q-learning”. In: *Machine Learning*. 279–292.